AthleteShield

Wearable-Based Sports Injury Risk Prediction

Sector: Sports Science & Athletic Performance Analytics
Document type: Project Brief, Technical Architecture & Literature Survey
Version 1.0 — August 2026

---

1. Executive Summary

AthleteShield is a wearable-sensor-driven machine learning system that continuously estimates an athlete's short-term risk of non-contact musculoskeletal injury — principally ACL tears, hamstring strains, and other lower-extremity soft-tissue injuries. It fuses inertial measurement unit (IMU) signals, GPS/GNSS-derived external training load, and heart-rate-variability (HRV)-derived internal load into a rolling, per-athlete feature store, and scores injury risk using tree-ensemble models (Random Forest, XGBoost) with an optional temporal deep-learning layer (LSTM / 1D-CNN) for sequence-sensitive patterns. The system is designed to run as an edge application on a team's existing wearable data hub, surfacing traffic-light risk flags to coaches and medical staff so that training load can be proactively adjusted before injury occurs, rather than after.

2. Problem Statement

Non-contact injuries are the single largest controllable driver of athlete availability loss in professional sport. ACL ruptures routinely end or derail careers, and hamstring strains are among the most frequently recurring injuries in football, athletics, and other sprint-based sports. Team medical and performance staff currently combine wellness questionnaires, GPS load reports, and subjective clinical judgment ("gut feel") to decide who trains and who rests. This approach does not scale, is not consistent across staff, and reacts to fatigue only after it becomes clinically apparent. The core problem AthleteShield addresses is: *can objective, continuously streamed biomechanical and physiological signals be used to flag elevated non-contact injury risk early enough for coaching staff to intervene with a load adjustment?*

2.1 Why this matters

- **Financial impact:** Elite clubs report multi-million-dollar wage costs tied up in injured players unable to compete, alongside transfer-market and match-outcome losses.
- **Career impact:** ACL and recurrent hamstring injuries are strongly associated with reduced career longevity and post-injury performance decrements.
- **Data availability gap:** Teams already own GPS vests, IMU pods, and HR straps generating dense time-series data that is mostly used for post-hoc load reporting rather than predictive modelling.
- **Reactive vs. proactive care:** Current practice intervenes after soreness or a wellness-score dip is reported; a predictive layer can shift intervention earlier in the fatigue-accumulation curve.

3. Proposed Approach

AthleteShield treats injury risk estimation as a supervised, time-windowed classification problem. Raw multi-sensor streams are aggregated per athlete into overlapping sliding windows (ranging from single-session to 28-day acute:chronic windows), from which sport-science-validated load features are engineered. These features are passed into ensemble tree models that output a calibrated injury-risk probability for each athlete, refreshed after every session.

3.1 Modelling pipeline stages

- **Ingestion:** Stream raw IMU (tri-axial accelerometer + gyroscope), GPS/GNSS positional data, and HRV/RR-interval data from wearable hubs via BLE/ANT+ to a local edge gateway.
- **Cleaning & synchronisation:** Denoise signals, interpolate short dropouts, and time-align multi-sensor streams to a common clock.
- **Windowing:** Apply sliding windows at multiple granularities — within-session (for biomechanical asymmetry/fatigue markers) and multi-day (for cumulative load, e.g. 7-day:28-day acute-to-chronic workload ratio, ACWR).
- **Feature engineering:** Derive PlayerLoad/Body-Load metrics, high-speed-running and sprint-distance counts, deceleration counts, jump/landing asymmetry indices from IMU, RMSSD/SDNN HRV trend features, and monotony/strain indices.
- **Model training:** Train Random Forest and XGBoost classifiers (primary), benchmarked against logistic regression and, optionally, an LSTM/temporal-CNN branch for raw-signal sequence patterns, using class-imbalance-aware techniques (SMOTE, class weighting) given the rarity of injury events.
- **Explainability:** Apply SHAP value analysis so medical staff can see which load factors are driving an individual athlete's risk score, not just the score itself.
- **Deployment:** Package the trained model as an edge-deployable inference service on the team's wearable hub or an on-premises server, refreshed with nightly batch retraining as new outcome data accrues.

4. System Architecture

The architecture is organised as a layered pipeline: sensing at the edge, ingestion and feature computation on a local hub (to minimise latency and keep raw biometric data on-premises for privacy), a model layer that ensembles tree-based and optional deep-sequence models, and a decision-support layer that turns probabilities into coach- and clinician-facing actions.

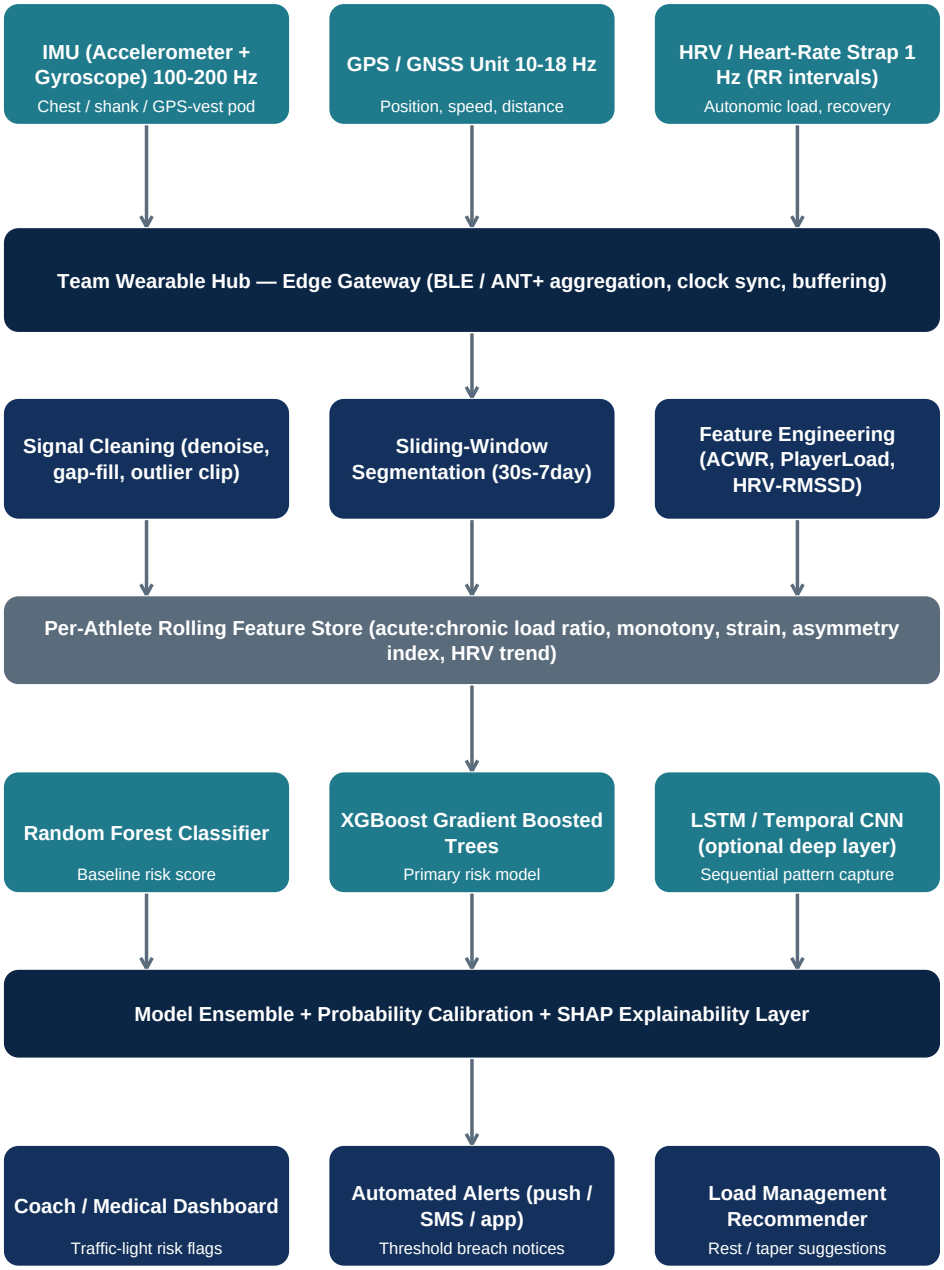

Outcome logging (confirmed injuries / near-misses) feeds back into the training set for periodic retraining

Figure 1. AthleteShield end-to-end architecture, from wearable sensing to coach-facing risk alerts, with an outcome-logging feedback loop for periodic retraining.

4.1 Layer-by-layer description

Layer	Function	Key Technology
Sensing	Capture raw IMU, GPS/GNSS, and HRV signals directly on the athlete's body.	Catapult/STATSports-class GPS vests, chest/shank IMU pods, chest-strap HR monitors
Edge gateway	Aggregate multi-athlete streams, synchronise clocks, buffer during connectivity loss.	BLE/ANT+ hub, local time-series buffer, on-prem server
Preprocessing	Denoise, interpolate gaps, and segment streams into sliding windows.	Butterworth/Kalman filtering, windowing (30 s to 28-day)
Feature store	Maintain a per-athlete rolling table of engineered load and physiological features.	ACWR, PlayerLoad, monotony/strain, HRV-RMSSD trend, asymmetry index
Model layer	Score injury-risk probability per athlete per session.	Random Forest, XGBoost, optional LSTM/1D-CNN branch
Ensemble & explainability	Calibrate probabilities and attribute risk drivers.	Stacked/weighted ensemble, SHAP values, Platt scaling
Decision support	Convert scores into actionable outputs for staff.	Coach dashboard, threshold alerts, load-taper recommender
Feedback loop	Log confirmed injuries/near-misses to improve future model versions.	Outcome tagging, scheduled retraining pipeline

4.2 Design considerations

- **Edge-first deployment:** Running inference on the team's local hub avoids sending raw biometric data off-site, reducing latency and easing data-governance concerns around athlete health data.
- **Class imbalance handling:** Injury events are rare relative to total training sessions (often <5% positive class), so the pipeline uses weighted loss functions, SMOTE-style oversampling, and recall/F1/AUC-oriented evaluation rather than raw accuracy.
- **Multi-horizon prediction:** Separate model heads or window configurations for next-session risk (1-2 days ahead) versus next-month risk (cumulative load-driven), since different features dominate at each horizon.
- **Explainability by design:** SHAP-based feature attribution is treated as a first-class output, not an afterthought, since medical staff must be able to justify a rest decision to players and management.
- **Human-in-the-loop:** The system flags elevated risk; it does not autonomously withdraw players from training or matches. Final load decisions remain with coaching and medical staff.

5. Datasets

No single public dataset currently combines match-grade GPS load, IMU biomechanics, HRV, and confirmed injury outcomes for professional athletes end-to-end, largely due to the sensitivity of athlete health records. AthleteShield's data strategy therefore combines injury-outcome datasets, wearable/GPS load datasets, and general-purpose IMU/human-activity-recognition (HAR) datasets that are adapted for gait and movement-quality feature extraction, supplemented by synthetic training-load data for pipeline development and stress-testing before deployment on real team data.

Dataset	Type / Content	Role in AthleteShield
Kaggle “Football Player Injury Data”	Player-level tabular records of match/training exposure, injury type, and recovery duration for footballers.	Primary source of labelled injury outcomes for supervised training and evaluation of the tree-ensemble classifiers.
Academic GPS/wearable sports-science cohort studies (e.g. youth soccer U10–U15 GNSS studies, professional-league GPS/training-load datasets published alongside peer-reviewed papers)	Session-level external load metrics: total distance, high-speed running, accelerations/decelerations, PlayerLoad, acute:chronic workload ratio, linked to injury/non-injury outcomes.	Provides realistic load-accumulation patterns and validated feature definitions (e.g. ACWR) for feature engineering and model benchmarking against published results.
PAMAP2 Physical Activity Monitoring Dataset	IMU (accelerometer, gyroscope, magnetometer) and heart-rate data from 3 body-worn units across 9 subjects performing 12+ activities including sport-relevant movements (e.g. running, cycling, soccer-like actions).	Adapted as a pretraining/transfer-learning source for IMU feature extractors (gait cadence, impact-peak detection, movement-symmetry features) before fine-tuning on team-specific biomechanics data.
MotionSense Dataset	Smartphone accelerometer/gyroscope/attitude data from 24 subjects performing walking, jogging, and stair activities at 50 Hz.	Used to validate lightweight on-device signal-processing and activity-segmentation logic before deployment on athlete-worn IMU pods.
Synthetic training-load datasets (sports-science-lab-generated)	Simulated multi-week training-load and injury-occurrence sequences generated to match published epidemiological injury-rate statistics.	Used for pipeline stress-testing, class-imbalance experimentation, and privacy-safe development/demo environments before a team's real data is available.

Note on data adaptation: PAMAP2 and MotionSense are general-purpose human-activity-recognition datasets, not injury datasets. They do not contain injury labels. Their role is restricted to (a) validating IMU signal processing, gait-segmentation, and movement-feature extraction code, and (b) transfer-learning pretraining for a movement-quality encoder, which is then fine-tuned on sport-specific, injury-labelled data such as the Kaggle football injury dataset and academic GPS-load cohorts. Using activity-recognition datasets as a direct substitute for injury-labelled data would not be methodologically valid.

6. Literature Survey

The table below summarises peer-reviewed and preprint literature (2023 onward) most directly relevant to AthleteShield's approach: wearable-sensor-driven, machine-learning-based sports injury risk prediction.

#	Title	Year	Venue	Relevance to AthleteShield
1	**Practice of wearable devices combined with deep learning algorithms in predicting athletic injury risk**	2025	Digital Health / SAGE	Reviews how deep-learning architectures paired with wearable IMU/GPS data are applied to injury-risk prediction, highlighting gains from temporal and multimodal fusion models over single-sensor approaches.
2	**A novel lower extremity non-contact injury risk prediction model based on multimodal fusion and interpretable machine learning**	2022–2023	Frontiers in Bioengineering & Biotechnology (PMC)	Proposes a fused biomechanical + physiological feature model with SHAP-based interpretability for lower-limb non-contact injury risk, directly informing AthleteShield's explainability layer.
3	**Predicting Injuries in Football Based on Data Collected from GPS-Based Wearable Sensors**	2023	Sensors, 23(3):1227	Uses Catapult GPS-vest data with an XGBoost classifier on an imbalanced injury/non-injury dataset, reporting 90% accuracy and 92% precision; a direct architectural precedent for AthleteShield's GPS-to-XGBoost pipeline.
4	**Artificial intelligence and wearable sensors in sports injury risk prediction: current status and future perspectives**	2026	Annals of Medicine (Taylor & Francis)	Narrative review covering IMU, EMG, and flexible-electronics wearables combined with deep-learning fusion models; identifies data heterogeneity, generalisability, and privacy as open challenges relevant to edge deployment.
5	**Construction and application of a model for predicting athletes' injury risk based on machine learning**	2024–2025	PeerJ / PMC	Builds a supervised ML pipeline on athlete training-load and physical-screening data, comparing ensemble tree methods against baseline regression for injury-risk classification.
6	**Machine learning approaches to injury risk prediction in sport: a scoping review with evidence synthesis**	2024	British Journal of Sports Medicine	Synthesises GPS, wellness-score, RPE, and screening-test-based ML injury studies across sports; finds tree-based ensembles (Random Forest, XGBoost) consistently among top performers, supporting AthleteShield's model choice.
7	**Machine learning framework for predicting athletic injuries and optimising performance**	2025	BMC Sports Science, Medicine and Rehabilitation	Trains Random Forest, XGBoost, and neural-network classifiers on workload/recovery/performance features for collegiate athletes; Random Forest achieves the strongest accuracy and ROC-AUC, reinforcing tree-ensemble baselines.
8	**Prediction of football injuries using GPS-based data in professional football players: a machine learning approach**	2025	Frontiers in Sports and Active Living	Applies XGBoost to acute:chronic workload ratio (7-day:28-day) features derived from GPS data in a professional league context, closely matching AthleteShield's proposed ACWR feature-engineering step.
9	**Time-to-injury forecasting in elite football: a survival-analysis approach with tree-ensemble baselines**	2026	arXiv preprint	Benchmarks XGBoost, Random Forest, and logistic regression against a deep survival model (DeepHit) for next-day injury-risk forecasting, using F1/recall/AUC-weighted evaluation to handle severe class imbalance — directly relevant to AthleteShield's evaluation strategy.

#	Title	Year	Venue	Relevance to AthleteShield
10	**Integrating multimodal AI technologies for sports injury prediction and rehabilitation: a systematic review**	2024–2025	Journal of Human Sport and Exercise	Presents a five-stage injury-management pipeline integrating sensor fusion, time-series classification, and predictive analytics, framing injury prediction as one stage within a broader monitor-predict-intervene-rehabilitate loop.
11	**The role of AI and wearables in injury prediction and performance optimisation**	2025	Premier Science Journal	Surveys wearable-driven AI systems for real-time athletic performance and injury-risk monitoring, discussing biomechanical analytics and workplace/industrial cross-over applications of the same sensor stack.
12	**Diagnostic applications of AI in sports: a comprehensive review of injury risk prediction methods**	2024	Conference proceedings (Education & Multimedia Technology)	Reviews AI-driven analytics for optimising training and treatment decisions in physical education and sport settings, including practical deployment considerations for coach-facing tools.

Publication years reflect the source journal's listed publication or online-first date at time of writing; several 2026-dated entries are early-2026 online-first releases. Full reference list with links is provided in Section 9.

7. Model Evaluation Strategy

Because confirmed injuries are rare events relative to total exposure time, raw classification accuracy is a poor and misleading metric. AthleteShield follows the evaluation conventions established in the literature surveyed above.

- **Primary metrics:** Recall (sensitivity) and F1-score on the injury-positive class, plus ROC-AUC and precision-recall AUC, since missing a true injury risk (false negative) is far costlier than a false alarm.
- **Class imbalance handling:** Stratified k-fold cross-validation, SMOTE/ADASYN oversampling or class-weighted loss during training, and a held-out, chronologically later test split to avoid leakage across a season.
- **Calibration:** Platt scaling or isotonic regression applied post-hoc so that a reported "70% risk" score is meaningfully close to a 70% empirical injury rate among similarly scored athletes.
- **Explainability validation:** SHAP feature-importance rankings cross-checked against sports-medicine domain knowledge (e.g. ACWR > 1.5 and HRV suppression should rank highly) before clinical trust is extended to the model.
- **Prospective validation:** Before full deployment, risk flags are logged silently for a full season without influencing selection decisions, so that flagged-vs-actual injury outcomes can be compared retrospectively.

8. Challenges & Limitations

- **Label scarcity and imbalance:** Injuries are rare (often single-digit percentages of exposures), which limits statistical power and increases overfitting risk for any single team's dataset.
- **Data heterogeneity:** Different wearable vendors (Catapult, STATSports, Polar, Whoop) use different sampling rates, proprietary load metrics, and calibration standards, complicating cross-team model transfer.
- **Confounded causality:** Elevated load metrics correlate with, but do not deterministically cause, injury; models must be interpreted as risk-stratification tools, not diagnostic instruments.
- **Privacy and governance:** Athlete biometric and health data is highly sensitive; edge-first, on-premises processing and strict access controls are required, along with athlete consent and applicable data-protection compliance.
- **Generalisability across sports and populations:** A model trained on adult professional footballers is unlikely to transfer directly to youth athletes or different sports without recalibration, as noted across the surveyed reviews.

9. Future Directions

- **Federated learning across clubs:** Train shared models across multiple teams without centralising raw athlete data, addressing both the label-scarcity and privacy challenges simultaneously.
- **Digital-twin athlete modelling:** Combine biomechanical simulation with sensor data to model individualised tissue-loading thresholds rather than population-level cut-offs.
- **Richer sensor fusion:** Incorporate EMG, force-plate, and flexible-electronics strain sensors as they become field-deployable, following the multimodal-fusion direction highlighted in recent reviews.
- **Closed-loop intervention testing:** Move from observational validation to randomised load-adjustment trials that measure whether acting on AthleteShield's flags actually reduces injury incidence.

10. Full Reference List

Sage Journals (2025). Practice of wearable devices combined with deep learning algorithms in predicting athletic injury risk. https://journals.sagepub.com/doi/10.1177/18724981251380391

PMC (2022/2023). A novel lower extremity non-contact injury risk prediction model based on multimodal fusion and interpretable machine learning. https://www.ncbi.nlm.nih.gov/pmc/articles/PMC9520324/

Pilka, T., Grzelak, B., Sadurska, A., Gorecki, T., & Dyczkowski, K. (2023). Predicting Injuries in Football Based on Data Collected from GPS-Based Wearable Sensors. Sensors, 23(3), 1227. https://www.ncbi.nlm.nih.gov/pmc/articles/PMC9919698/

Taylor & Francis Online (2026). Artificial intelligence and wearable sensors in sports injury risk prediction: current status and future perspectives. Annals of Medicine. https://www.tandfonline.com/doi/full/10.1080/07853890.2026.2658879

PMC (2024/2025). Construction and application of a model for predicting athletes' injury risk based on machine learning. https://www.ncbi.nlm.nih.gov/pmc/articles/PMC12849558/

Br J Sports Med (2024). Machine learning approaches to injury risk prediction in sport: a scoping review with evidence synthesis. https://pmc.ncbi.nlm.nih.gov/articles/PMC12013557/

BMC Sports Science, Medicine and Rehabilitation (2025). Machine learning framework for predicting athletic injuries and optimising performance. https://link.springer.com/article/10.1186/s13102-025-01502-x

Frontiers in Sports and Active Living (2025). Prediction of football injuries using GPS-based data in Iranian professional football players: a machine learning approach. https://www.frontiersin.org/journals/sports-and-active-living/articles/10.3389/fspor.2025.1425180/full

arXiv (2026). Time-to-Injury Forecasting in Elite Female Football: A DeepHit Survival Approach. https://arxiv.org/pdf/2601.19479

Journal of Human Sport and Exercise (2024/2025). Integrating multimodal AI technologies for sports injury prediction and rehabilitation: Systematic review. https://www.jhse.es/index.php/jhse/article/view/integrating-multimodal-ai-technologies-sports-injury-prediction

Premier Science (2025). The Role of AI and Wearables in Injury Prediction and Performance Optimization. https://premierscience.com/pjai-25-971/

Premier Science (2025). Machine Learning Approaches to Injury Risk Prediction in Sport. https://premierscience.com/pjcs-25-918/

Reiss, A., & Stricker, D. (2012). PAMAP2 Physical Activity Monitoring Dataset. UCI Machine Learning Repository.

Malekzadeh, M. et al. (2018/2019). MotionSense Dataset for Human Activity and Attribute Recognition. IoTDI '19 / PMC.

Kaggle. Football Player Injury Data (community dataset).

Prepared as a project brief for AthleteShield. All cited works are publicly available academic or preprint sources; links reflect the source pages retrieved at time of writing and should be re-verified before formal submission.